

Lecture 4:

Basics of Parallel Programming

Beakal Gizachew Assefa

So far . . .

- Importance of Parallel Programming
- Parallel Architectures
- Scalability, Speedup, and Performance
- Performance Modeling
 - With a pencil-and-paper
 - Will resume with case studies later
- Today
 - Let's write our first parallel program
 - Reading Chapter 2.4.1–2.4.3

How to Parallelize?

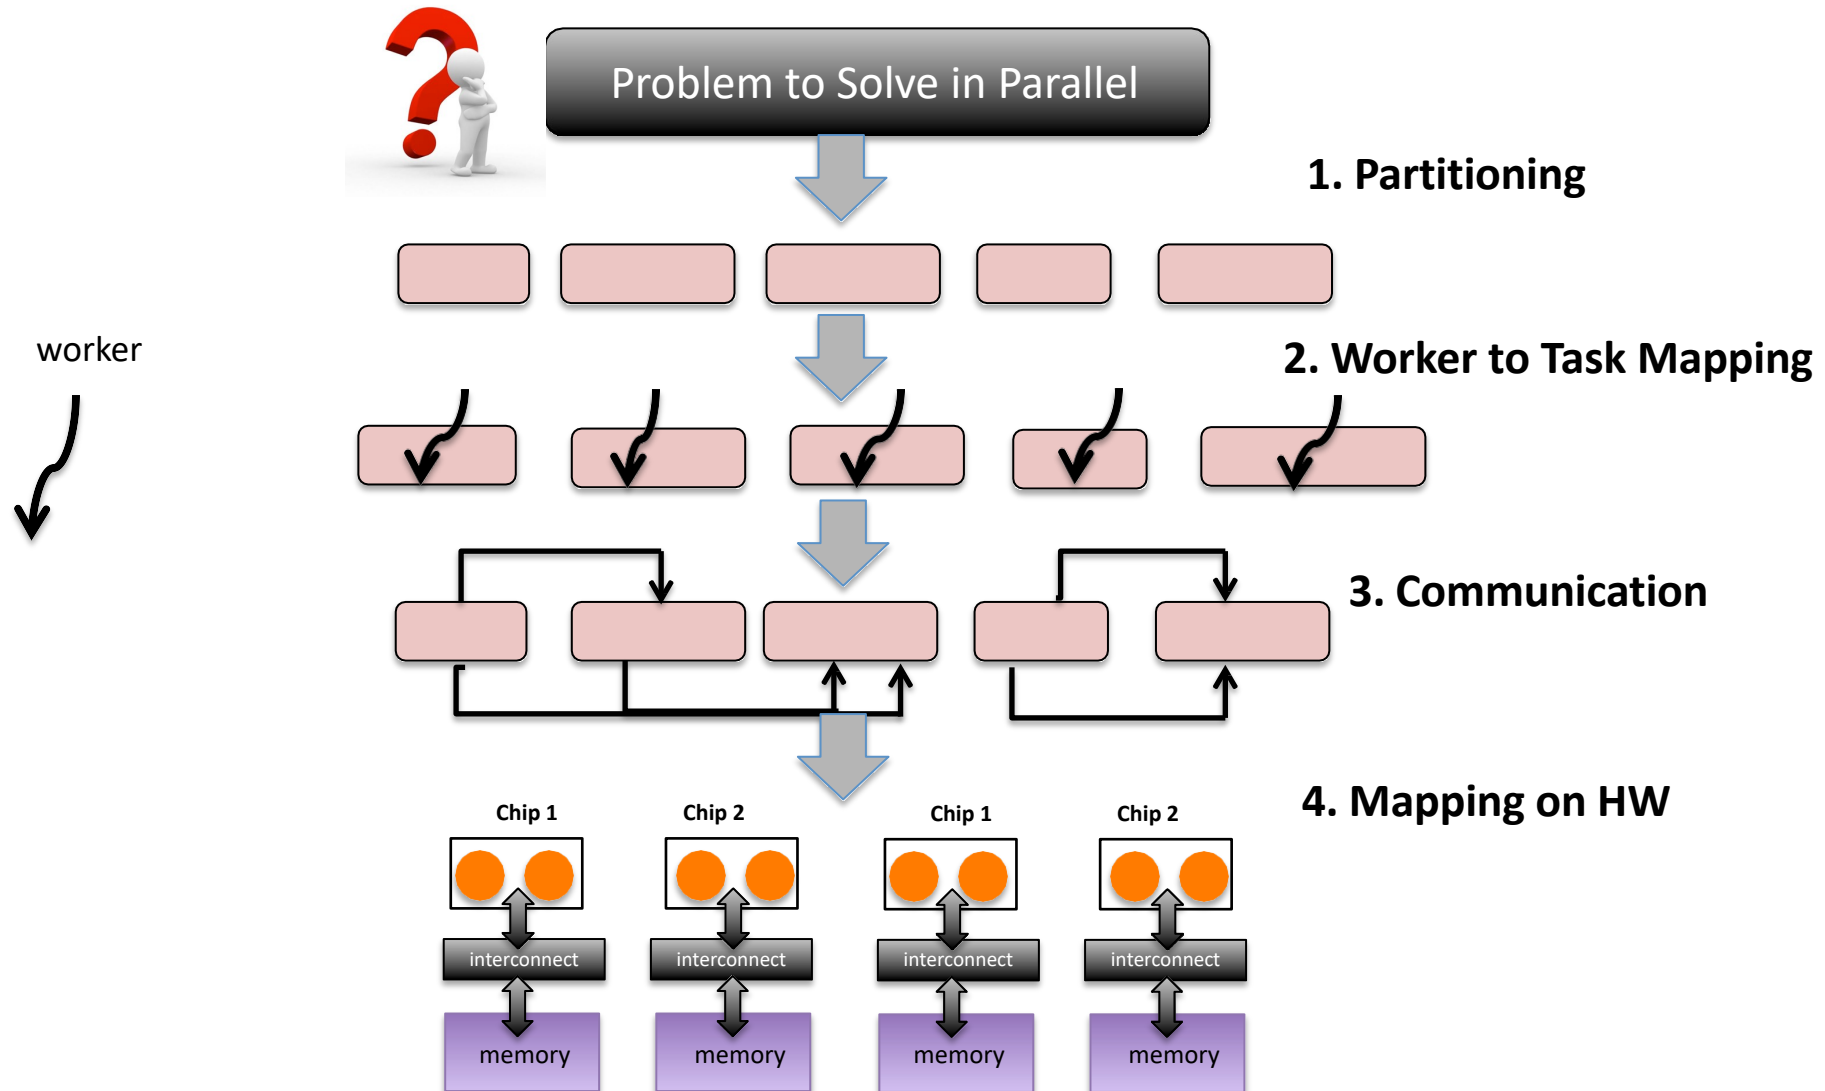
- Assume we are starting with a **sequential algorithm** and trying to modify it to execute in parallel
 - Not always the best strategy, as sometimes the best parallel algorithms are NOTHING like their sequential counterparts
 - But useful since you are accustomed to sequential algorithms

How to Parallelize?

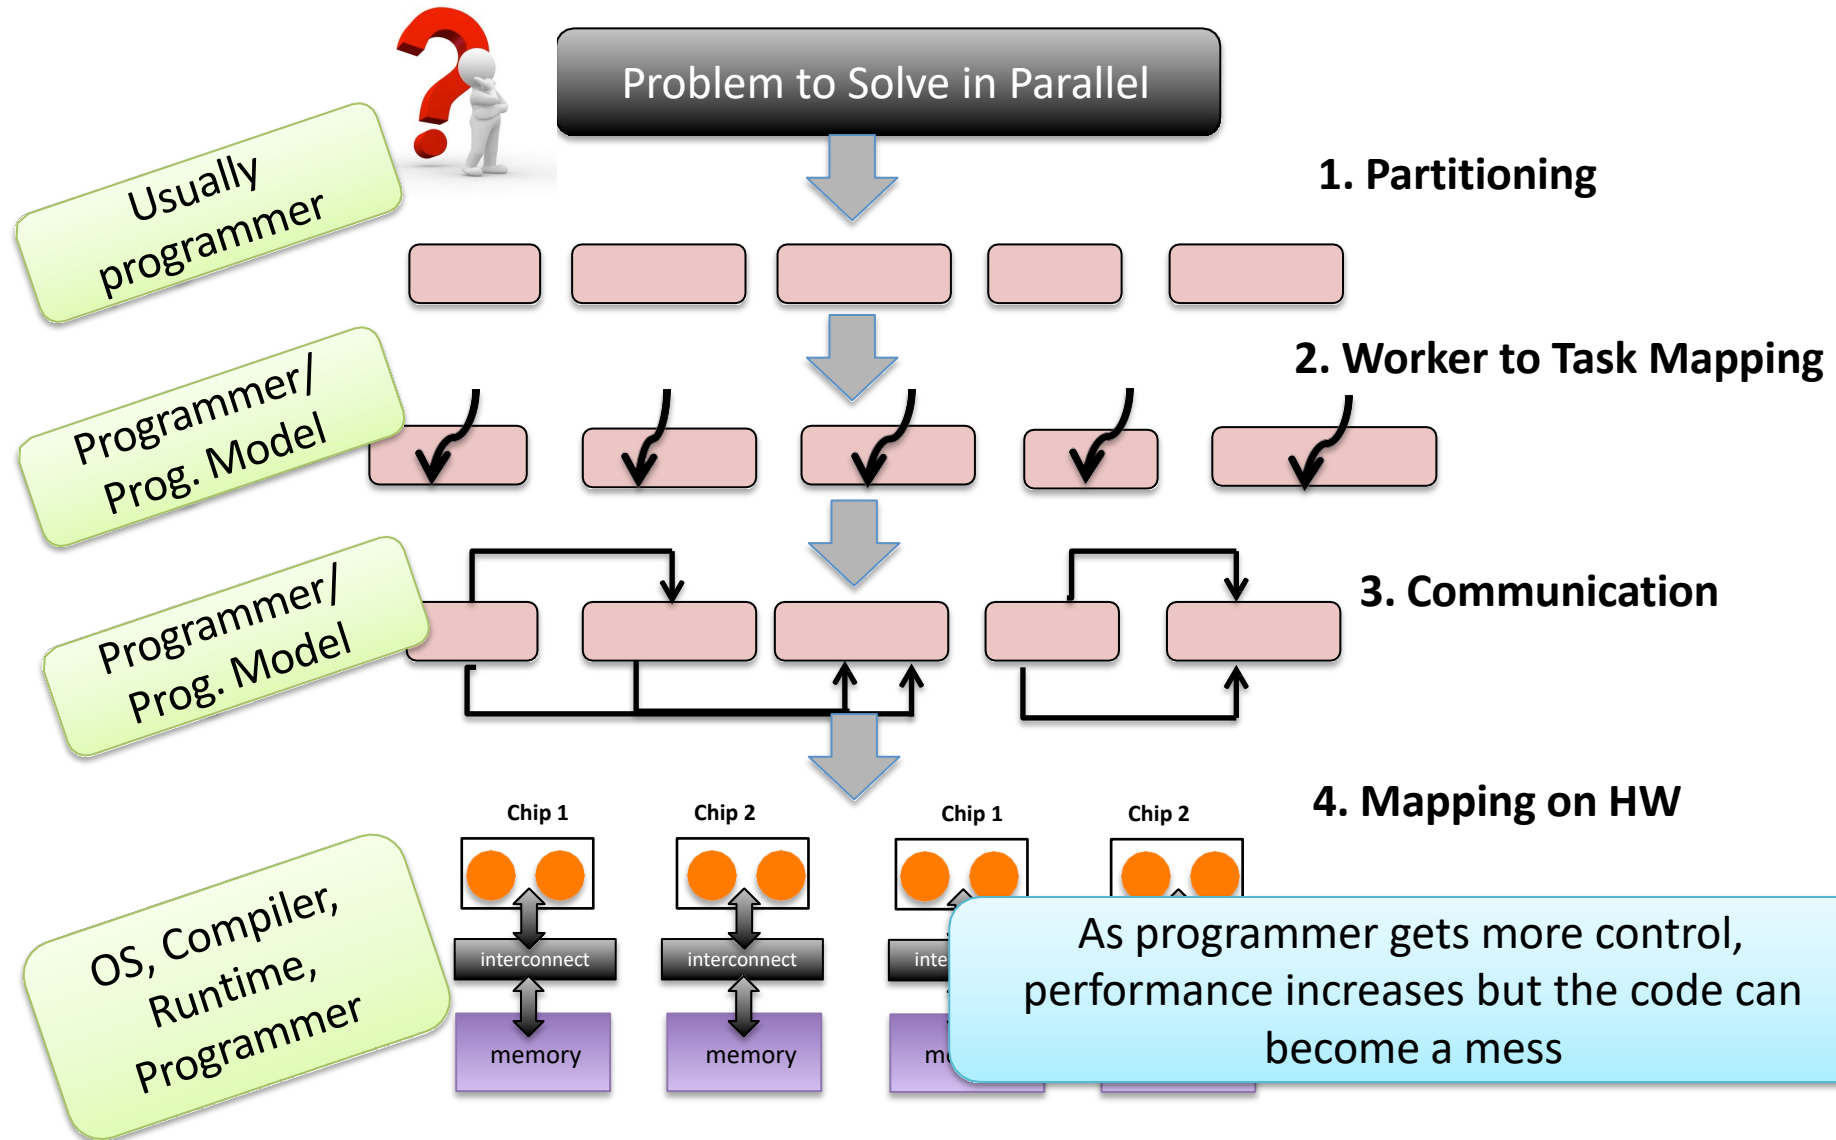
- Identify work that can be performed in parallel
 - Remember Amdahl's Law: reduce the serial fraction as much as you can for scalability
- Partition work among workers
- Manage data access, communication and coordination of workers
- Remember our ultimate goal is “**Performance**”
 - Ignore architectural details at first
 - Then, optimize for architecture
 - What is the downside of this?

Performance Portable is
more desirable

Reasoning about a Parallel Program



Who is Responsible?



Parallelization Steps

- 1. Partitioning: Divide the computation into tasks
 - Will cover throughout the semester
- 2. Worker-to-Task Mapping: Assign threads/workers to tasks to execute
 - Advanced topic
 - This can be very complicated when dealing with irregular applications (e.g. graphs)
- 3. Communication: determine what communication needs to be carried out among the tasks
 - Will cover throughout the semester
- 4. Mapping: Mapping threads to hardware execution units
 - Advanced topic
 - Will cover with respect to NUMA--architectures and GPUs

1.Partitioning


- **Computation Partitioning:**
 - Divide the sequential computation among parallel threads/processors/computations
 - The focus here should be on identifying **tasks** that can be executed in parallel.
- **Data Partitioning**
 - Also known as data decomposition or domain decomposition
 - Partition the data operated on by the computation
- One might imply the other
 - Think of whichever is easier, then think the corresponding partitioning

When partitioning the computation

- There are two important aspects:
 - Granularity of tasks
 - Size of each task
 - Static vs dynamic partitioning
 - New tasks can be discovered as program runs
 - Preserving data dependencies
 - Keeping the data values consistent with respect to the sequential execution.

Granularity of Tasks

- Fine--grained



- Tasks are small
- Increases the degree of parallelism 😊
- Increases the parallelization overhead ☹️

- Coarse--grained



- Tasks are coarser
- Decreases the degree of parallelism 😊
- Reduces the overhead 😊

- Finding the right granularity is the key


- Create at least enough tasks to keep all execution units on a machine busy
- Each task working--set should fit into L1 if not L2 cache

Data Dependencies

- One of the difficulties of parallel programming comes from the data dependencies between tasks
- Parallel execution has to obey the data dependencies otherwise we will end up with an incorrect program
- A formal definition:
 - A *data dependence* is an ordering on a pair of memory operations that must be preserved to maintain correctness.

A Simple Example

- On an N-by-N image, consider a computation that converts color to grayscale
 - Each color pixel is described by a triple (R, G, B) of intensities for red, green, and blue
 - Average method simply averages the values: $(R + G + B) / 3$ on each pixel
- Here, computation on each pixel is **independent**, no data dependencies between tasks
 - These types of parallelization are called ‘**embarrassingly parallel**’ algorithms



Another Example: Parallel Sum


- Compute n values and add them together
- Serial formulation:

```
sum = 0;
for (I = 0; I < N ; I++ )
{
    x= compute_next_value( . . . );
    sum  = sum + x;
}
```

- Parallel formulation?

Version 1: Naïve


- Computation Partitioning
 - Suppose each task computes a partial sum on n/t consecutive elements (t is the number of tasks)
 - Example: $n = 24$ and $t = 8$ tasks



- Workers to Task Mapping
 - Assume we have 8 cores/processors
 - Each worker/thread gets a task
 - Need to calculate the start index for each thread

Version 1: Naïve

- Example: $n = 24$ and $t = 8$ tasks (threads)



```
int items_per_task = n/t;  
int start = thread_id * items_per_task;  
  
for (i=start; i<start + items_per_task; i++) {  
    x = Compute_next_value(...);  
    sum += x;  
}
```

Correct?

Data Dependencies?


- Load/increment/store must be done *atomically* to preserve sequential meaning
 - More than one thread may update sum at the same time
- A *race condition* exists when the result of an execution depends on the *timing* of two or more events.
- **Mutual exclusion**: at most one thread can execute the code at any time

```
int items_per_task = n/t;
int start = thread_id * items_per_task;

for (i=start; i<start + items_per_task; i++) {
    x = Compute_next_value(...);
    sum += x;
}
```


Race Condition


- The value of sum is non--deterministic



Thread 1	Thread 2		Integer value
			0
read value		←	0
increase value			0
write back		→	1
	read value	←	1
	increase value		1
	write back	→	2

Race Condition

- The value of sum is non--deterministic



Thread 1	Thread 2		Integer value
			0
read value		←	0
	read value	←	0
increase value			0
	increase value		0
write back		→	1
	write back	→	1

Version 2: Add Locks

- Insert mutual exclusion (mutex) so that only one thread at a time is loading/incrementing/storing sum **atomically**
 - **Atomicity**: a set of operations is atomic if either they all execute or none executes. Thus, there is no way to see the results of a partial execution.

```
int items_per_task = n/t;
mutex m;
int start = thread_id * items_per_task;

for (i=start; i<start + items_per_task; i++) {
    my_x = Compute_next_value(...);
    mutex_lock(m);
    sum += my_x;
    mutex_unlock(m);
}
```

Now, it is correct!

Version 3: Reduce the use of Locks

- Acquiring lock brings overhead because it serializes parallel execution
- Lock only to update final sum from thread--private copy


```
int items_per_task = n/t;
mutex m;
int my_sum;
int start = thread_id * items_per_task;

for (i=start; i<start + items_per_task; i++) {
    my_x = Compute_next_value(...);
    my_sum += my_x;
}
mutex_lock(m);
sum+= my_sum;
mutex_unlock(m);
```

Can we eliminate the lock completely?

Version 4: Eliminate lock

- One of the threads can accumulate result
- Local sum is indexed by thread ID



Version 4: Eliminate lock

- One of the threads can accumulate result

```
int items_per_task = n/t;

shared int my_sum[t]; //number of threads
int start = thread_id * items_per_task;

for (i=start; i<start + items_per_task; i++) {
    my_x = Compute_next_value(...);
    my_sum[thread_id] += my_x;
}

if (thread_id == 0 ) //master thread
{
    sum = my_sum[0];
    for(i=1; i< t; i++) sum+ = my_sum[i];
}
```




Correct?

Synchronization: Barriers

- Sum is incorrect if `master` thread begins accumulating final result before other threads are done
- **Synchronization** is used to sequence control among threads or to sequence accesses to data in parallel code.
- How can we force the master to wait until the threads are ready?
 - A **barrier** is used to block threads from proceeding beyond a program point until all of the participating threads has reached the barrier.

Version 5: Add a barrier

- Ensure all the local sums are ready (all the threads are done calculating their local sums)



Version 5: Add a barrier

- Master waits for others to finish

Now it is
correct!

```
int items_per_task = n/t;
```

```
shared int my_sum[t]; //number of threads  
int start = thread_id * items_per_task;
```


```
for (i=start; i<start + items_per_task; i++) {  
    my_x = Compute_next_value(...);  
    my_sum[thread_id] += my_x;  
}
```

synchronize_threads(); // barrier for all participating threads

```
if (thread_id == 0 ) //master thread  
{  
    sum = my_sum[0];  
    for(i=1; i< t; i++) sum+ = my_sum[i];  
}
```

Version 6: Improve Performance

- Now, our implementation is correct, we can try to improve its performance




Version 6: Tree Sum

- Threads form a tree to accumulate sum
- Sum is calculated in $\log(t)$ steps, where t is number threads/processors
- For large t , it makes a performance difference
- For example, $N=1M$, $t = 1000$
 - Each worker computes N/t elements: $1M/1000 = 1000$ elements, then we have 1000 partial sums
 - If only master computes,
 - We have 1000 adds by master (serialization)
 - Total time= Time(partial sum) + MasterTime(global sum)
 - In tree sum
 - Total time= Time(partial sum) + $\log(t)$
 - We have only 10 adds for partial sums

Version 6: Tree Sum

- Need to add synchronization points (not necessarily a global barrier)



Data Dependencies?

- Dependence on sum across iterations/threads?
 - Reordering ok since operations on sum are associative
- Calculating
 - $(((((1+4)+3)+9)+2)+8)$ is the same as
 - $(1 + 4 + 3) + (9 + 2 + 8)$



- May get slightly different results on floating point operations
 - because of rounding in hardware
 - Real numbers are approximated in hardware

Lessons Learnt from Parallel Sum

- The sum computation had a race condition or data dependence.
- We used mutex and barrier synchronization to guarantee correct execution.
- We performed mostly local computation to increase parallelism granularity across threads.
- What were the overheads we saw with this example?
 - Extra code to determine portion of computation
 - Locking overhead: inherent cost plus contention
 - Load imbalance: use tree sum

Acknowledgments

- These slides are inspired and partly adapted from
 - Mary Hall (Univ. of Utah)
 - The course book (Pacheco)