

# **Lecture 8:**

# **Task Parallelism**

**Beakal Gizachew Assefa**

# Today

- Types of Parallelism
  - Data Parallelism
  - Task Parallelism
- Task Parallelism in OpenMP

# OpenMP Summary

- Work sharing
  - parallel, parallel for
  - scheduling directives: `static()`, `dynamic()`, `guided()`
- Data sharing and variable scopes
  - shared, private, reduction
- Environment variables
  - **OMP\_NUM\_THREADS**
- Library
  - E.g., `omp_get_num_threads()`, `omp_get_thread_num()`

# Data Dependence

- A ***data dependence*** is an ordering on a pair of memory operations that must be preserved to maintain correctness.
- **Question:** When is parallelization guaranteed to be safe?
- **Answer:** If there are no data dependences across reordered computations.
- **Definition:** Two memory accesses are involved in a data dependence if they may refer to the same memory location and one of the accesses is a **write**.

# Preserve Dependences

- Fundamental Theorem of Dependence:
  - Any reordering transformation that preserves every dependence in a program preserves the meaning of that program.
- Parallelization
  - Computations that execute in parallel between synchronization points are potentially reordered.
  - Is that reordering safe? According to our definition, it is safe if it preserves the dependences in the code.

```
#pragma omp barrier  
  
. . .  
#pragma omp for nowait  
for()  
  
. . .  
#pragma omp barrier
```

**It's programmer's responsibility to ensure there is no violation of data dependences.**

# Parallel Regions

```
#pragma omp parallel
{

}
#pragma omp parallel for
for (i=0; i<size; i++) {

}
#pragma omp parallel {

}
```

```
#pragma omp parallel
{

}
```

- Avoid creating too many parallel regions if you can
  - Each time, threads are created and joined!!!
- Avoid also expanding a parallel region into multiple files
  - Find a compromise

# Parallel Loops

- Guarantee that the same thread will compute the same code range if the same scheduling is used in the two loops

Can recommend because it can only introduce a performance bug even if the loop schedulings are different  
First loop performs the first-touch.

```
#pragma omp parallel
{
    #pragma omp for schedule(static) nowait
    for()
        a(i)=

    #pragma omp for schedule(static)
    for()
        . . = a(i)
}
```

I don't recommend to use it because it can introduce a bug if you decide to change the scheduling type  
Fuse loops if possible.

```
#pragma omp parallel for
for (i=0; i<N; i++)
    A[i] = 1.0;

. . .
#pragma omp parallel for
for (i=0; i<N; i++)
    //calculation on A
}
```

# Conditional Parallelization

```
if (scalar expression)
```

***Only execute in parallel if expression evaluates to true  
Otherwise, execute serially***

```
#pragma omp parallel if (n > threshold) \  
shared(n,x,y) private(i)  
{  
    #pragma omp for  
    for (i=0; i<n; i++)  
        x[i] += y[i];  
} /*-- End of parallel region --*/
```



# Single and Master Constructs

- Only one thread in team executes code enclosed
- Useful for things like I/O or initialization
- **Implicit barrier on exit**

```
#pragma omp single {  
    <code-block>  
}
```

- Similarly, only master executes code
- **No implicit barrier on exit**

```
#pragma omp master {  
    <code-block>  
}
```

# Types of Parallelism

- Task parallelism
  - Partition various distinct tasks carried out solving the problem in parallel.
- Data parallelism
  - Partition the data used in solving the problem in parallel
  - Each thread/process carries out similar operations on it's part of the data.


# Task vs Data Parallelism

- Assume there are 30 students
  - Thus 30 exams and
  - 3 questions in the exam




# Grading


- Ruth, Meba and Ikbal are grading the exams



10 exams



10 exams






10 exams

- Scenario #1: Each of us gets 10 students to grade

# Grading

- Ruth, Meba and Ikbal are grading the exams



- Scenario #2: Each of us grades only 1 question

# Task vs Data Parallelism

- Scenario #1: Data parallelism
  - Data (exam papers) is divided among all, the same work is performed on the data
- Scenario #2: Task parallelism
  - Tasks (questions) are divided, each TA gets a different question

# Task vs Data Parallelism-- Discussion



- (a) Identify a portion of home building that can employ *data parallelism*, where “data” in this context is any object used as an input to the home--building process, as opposed to tools that can be thought of as processing resources.
- (b) Identify *task parallelism* in home building by defining a set of tasks. Work out a schedule that shows when the various tasks can be performed.
- (c) Describe how task and data parallelism can be combined in building a home. What computations can be reassigned to different workers to balance the load?

# Task Parallelism

- OpenMP supports for task parallelism
  - Parallel sections: different threads execute different code
    - OpenMP Pre-3.0
  - Tasks (NEW): tasks are created and executed at separate times
    - OpenMP 3.0




# OpenMP Sections



# Parallel Sections in OpenMP

```
#pragma omp parallel shared(n,a,b,c,d) private(i)
{
    #pragma omp sections
    {
        #pragma omp section
        for (i=0; i<n; i++)
            d[i] = 1.0/c[i];
        #pragma omp section
        for (i=0; i<n-1; i++)
            b[i] = (a[i] + a[i+1])
    } /*-- End of sections --*/
} /*-- End of parallel region
```



# OpenMP Sections

- The SECTIONS directive is a work--sharing construct
  - It specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- Independent SECTION directives are nested within a SECTIONS directive.
- Each SECTION is executed once by a thread in the team.
- Different sections may be executed by different threads.
  - It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.

# Tasks in OpenMP 3.0

- A task has
  - Code to execute
  - A data environment (shared, private, reduction)
  - An assigned thread that executes the code and uses the data
- Two activities: packaging and execution
  - Each encountering thread packages a new instance of a task
  - Some thread in the team executes the thread at a later time

# Definitions

- Task construct – task directive plus structured block

```
#pragma omp task
{
    //structured block
}
```

- Task – the package of code and instructions for allocating data created when a thread encounters a task construct
- Task region – the dynamic sequence of instructions produced by the execution of task by a thread

# When/Where are tasks complete?

- At thread barriers, explicit or implicit
  - Applies to all tasks generated in the current parallel region
- At task barriers
  - Wait until **all tasks defined in the current task** have completed
  - `#pragma omp taskwait`
  - Note: applies only to tasks generated in the current task

# Example

- Parallel pointer chasing using tasks

```
#pragma omp parallel
{
    #pragma omp single private(p)
    {
        p = listhead;
        while (p) {
            #pragma omp task firstprivate(p)
            process(p)
            p = next (p)
        }
    }
}
```

Task gets its own copy of p  
initialized to the value of p when  
the task is defined

for each p in the list, create a  
new task to process it

# Example

- Parallel pointer chasing on multiple lists using tasks


```
#pragma omp parallel
{
    #pragma omp for private(p)
    for (int i=0; i < numList; i++)
        p = listheads[i];
        while (p){
            #pragma omp task
            process(p)
            p = next (p)
        }
    }
}
```



# Example

- Postorder tree traversal
- Parent task suspended until children tasks complete (A, C, E, D, B, H, I, G, F)

```
void postorder (node* p){  
  
    if (p-> left)  
        #pragma omp task  
        postorder(p-> left);  
    if (p-> right)  
        #pragma omp task  
        postorder (p->right);  
  
    //wait for descendants  
    #pragma omp taskwait  
    process (p-> data);  
}
```



# Final Remarks on Tasks

- Integrated into OpenMP 3.0
- Flexible model for irregular parallelism
  - Graph traversal
- Nested parallelism where tasks creates new tasks

```
#pragma omp parallel
{
    if(tid == 0)
        //do this
    else
        //do that
    #pragma omp barrier
}
```

If you have a work sharing construct inside  
else, then tid=0 will not reach that construct!  
Result is incorrect

```
else {
    #pragma omp for
    for ... {
        ...
    }
}
```

# Acknowledgments

- These slides are inspired and partly adapted from
  - Mary Hall (Univ. of Utah)
  - The course book (Pacheco)
  - <https://computing.llnl.gov/tutorials/openMP/>